Recent text-to-motion (T2M) methods learn structured motion distributions conditioned on text or spatial constraints (Petrovich et al., 2022; Tevet et al., 2023; Chen et al., 2023; Zhang et al., 2023a; Guo et al., 2024; Xie et al., 2024).

Building on these advances, **Kimodo-Policy** combines a pretrained T2M backbone with three-stage root–body–hand generation for humanoid VLA control. ❶ **Strong Whole-Body Motion Priors.** We transfer Kimodo's pretrained inter-joint and temporal priors to high-dimensional, visually grounded control (Fig. 2(c)) (Rempe et al., 2026), easing adaptation from limited task-specific humanoid demonstrations. ❷ **Strong Component-Specific Motion Modeling.** Dedicated generators (Fig. 2(f)) predict root translation and heading, root-conditioned body motion, and hand controls conditioned on future body features and visual observations. Their dependencies support component specialization and whole-body coordination. We pretrain on diverse humanoid demonstrations to generate actions conditioned on visual observations, language instructions, and motion history. With 678M total and only 92M trainable parameters, Kimodo-Policy shows consistent gains from more VLA pretraining data, larger models, and stronger T2M priors. Our main contributions are: